

VLM Classification Experiment: LLaVA-OneVision (Ordered vs Shuffled)

Model: LLaVA-OneVision-Qwen2-0.5B | Frames per video: 4 | Videos evaluated: 25

Condition	Accuracy
Ordered Frames	40.00%
Shuffled Frames	44.00%
Drop (Ordered - Shuffled)	-4.00pp

Case Study 1: Pushing [something] from left to right

Video ID: 85446 | Ground Truth: 0 (Pushing [something] from left to right)



Frame 1



Frame 2



Frame 3



Frame 4


Condition	Predicted Label	Correct?
Ordered	0	Yes
Shuffled	0	Yes

Case Study 2: Pushing [something] from left to right

Video ID: 91241 | Ground Truth: 0 (Pushing [something] from left to right)



Frame 1



Frame 2



Frame 3



Frame 4


Condition	Predicted Label	Correct?
Ordered	1	No
Shuffled	0	Yes

Case Study 3: Pushing [something] from left to right


Video ID: 201692 | Ground Truth: 0 (Pushing [something] from left to right)



Frame 1



Frame 2



Frame 3



Frame 4

Condition	Predicted Label	Correct?
Ordered	0	Yes
Shuffled	2	No